

Predicting Distribution of Dublin Bikes

Con O’Leary; olaoghac@tcd.ie

August 14, 2021

1 Evaluation

I believe R^2 evaluation is appropriate for measuring success in this task. A prediction that is off by one bike is probably only going to be a fifth as much of an issue (to a hypothetical user availing of the model) as a prediction that is off by five bikes. As such, that R^2 evaluation measures how correct the variance of predictions are—and that variance accounts equally for the frequency of errors and the magnitude of errors—means that error in the model is represented proportionally to how much of an issue is presented to the objective of the task.

2 Data procurement

There are considerable time gaps in *Dublinbikes 2020 Q1 usage data*. I recognised this accidentally, when making a provisional graph that showed the percent bike occupancy of each station throughout the full duration of the data. The sections of the x-axis that are atopped by red show the final station to be plotted (plotted in red) fluctuating in percent occupancy; they show where there is data. Where there is a mesh of colours, is where the graphing library is adjoining lines across a data-less span. **The end of the second data-less span is the 27th of January 2020, and so I processed the information from then onwards.** The reason why I was adamant to exclude any dates before the 27th that did have full data is so that I would not have more data for certain days of the week than I have for others: I was worried this might lead to the models orientating themselves more around certain days.

3 Baseline Approach

I first went about a linear regression. Noting the cyclical nature of the amount of bikes in a station, I derived polynomial features from the amount of bikes in a given station over the total time. I disabled the shuffle parameter in my train-test split, as it seemed testing would be improper if, for example, we were looking for a prediction for the amount of bikes in station x in 30 minutes time when the prior training had included the amount of bikes in station x in 25




Figure 1: % bike occupancy for every station over the total duration

and 35 minutes time. Understandably, the x-axis of the training data spanning months and the x-axis of the testing data spanning weeks was not conducive at all, with any degree of polynomial features.

Next, with the same test-train division, I went about training 7 linear regressions—one for each day of the week. I trained the regression for each weekday on their respective polynomial features, derived from the average amount of bikes in a given station, on the respective weekday, at a given 5-minute interval. When this approach fell through I didn't stick around to figure out why, as I was looking for a baseless and not a complex model.

Finally, I abandoned the use of models and established a baseline directly from the average bike occupancies of the days of the week.

The results of which were solid; scoring near-positive R^2 scores. Looking through a comparison of predicted and desired y-values I thought some form of regularisation would increase the accuracy. To regularise, I diluted my predictions with the mean y-value for the respective weekday of the given prediction. Optimising the coefficient that determines the degree to which the original predictions are diluted, I scored a marginally positive score for one of my stations, and a near-positive for the other. Interestingly this was my only approach that worked better for the more central, less residential Custom House Quay station. Seeing as my baseline scored an R^2 score of roughly zero, this, as far as I understand, means its squared sum error was roughly that of the mean line of the training data. To gauge the efficacy of my baseline, as well as to see the effect of the regularisation, I plotted the R^2 scores of my baseline (with a range of values for my regularisation coefficient) versus the R^2 scores of the per-station mean of the training data. Realising the latter performed approximately as well as the approach I had developed, I adopted it as my baseline.

4 Approach 1

The idea of my first approach was to identify patterns between bikes disappearing in stations and bikes turning up in stations, and vice versa. As such I conceived a number of features: Chief among these is *bikes_changes_pastx*, which details for a given station at a given 5-minute interval the change in the number of bikes in the past x minutes. So that correlations between the changes detailed in *bikes_changes_pastx* could be found, a feature(s) representing time needed to be added. The particular components of time that seemed to matter for the purposes of representing these changes seemed to be the hour of day and the day of week. Figuring that an MLPRegressor might be a good means to represent these correlations, I made 7 inputs to represent the days of the week, and 23 to represent the time of the day. The 23 are represented as values between 0 and 1, such that 6,15 am would be represented by the 6th input having a value of 0.75, the 7th input having a value of 0.25, and the rest of the 23 inputs having a value of 0. Unlike this first time-related feature, I did not have the 7 inputs representing the days of the week transition gradually, but rather be 0 or 1, because the degree of transition between the days is already represented




Figure 2: Optimising the regularisation of baseline predictions

by the first time-related feature. Last of which, I included a feature that is the percent fullness of the station. Any features that did not already have values that ranged between 0 and 1 or -1 and 1 were normalised.

While I had hoped the model would primarily derive its correlations from the various components of *bikes.changes.pastx*, instances of the model that did not have a whole kitchen sink of *bikes.changes.pastx* components thrown at them outperformed those that did. In particular, (1) the inclusion of components pertaining to the change in amount of bikes over a duration greater than 10 minutes had no positive impact, and the inclusion of those greater than about 20 had a negative impact, and (2) the inclusion of components pertaining to the change in the amount of bikes in all stations had negative impacts. This was disappointing as I had hoped that a model having access to features that had the potential to describe inter-station bike transition patterns would discover such. Perhaps giving the changes in bike occupancy in all stations entailed too much irrelevant information. The way I would further this model if I had more time would be to only look at bike changes in the past x minutes for stations that can be travelled to and from in x minutes.

5 Approach 2

My second approach, at least as conceived in my mind, orientates around finding recurring behaviour in a station's bike occupancy changes. To identify these




Figure 3: % bike occupancy for every station over the total duration

recurrences I fit a KNN algorithm with features, such that the neighbours of a given datapoint for a given station are the datapoints for that station where the conditions of the features most closely resemble that of the datapoint in question. Like the first approach, it uses components from *bikes_changes_pastx*, as well as the stations percent occupancy. Time is represented differently to first approach. If the hour of day were to be represented as in the first approach, the KNN algorithm would regard 01:00 and 02:00 to be neighbourly, although it would not regard 23:00 and 00:00 to be neighbourly. As such, I represented the time of day as two different values, which are the xs and ys of a geometric circle. This means 23:00 and 00:00 would be regarded by the algorithm as being as neighbourly as 01:00 and 02:00.

I think this approach took poorly to the components of *bikes_changes_pastx* pertaining to more than 10 minutes ago because, unlike the multi-layer perceptron model, the KNN algorithm does not weight coefficients for its features, so components pertaining to greater durations have the same weighting as the more-relevant *past5* and *past10* components.

6 Conclusion

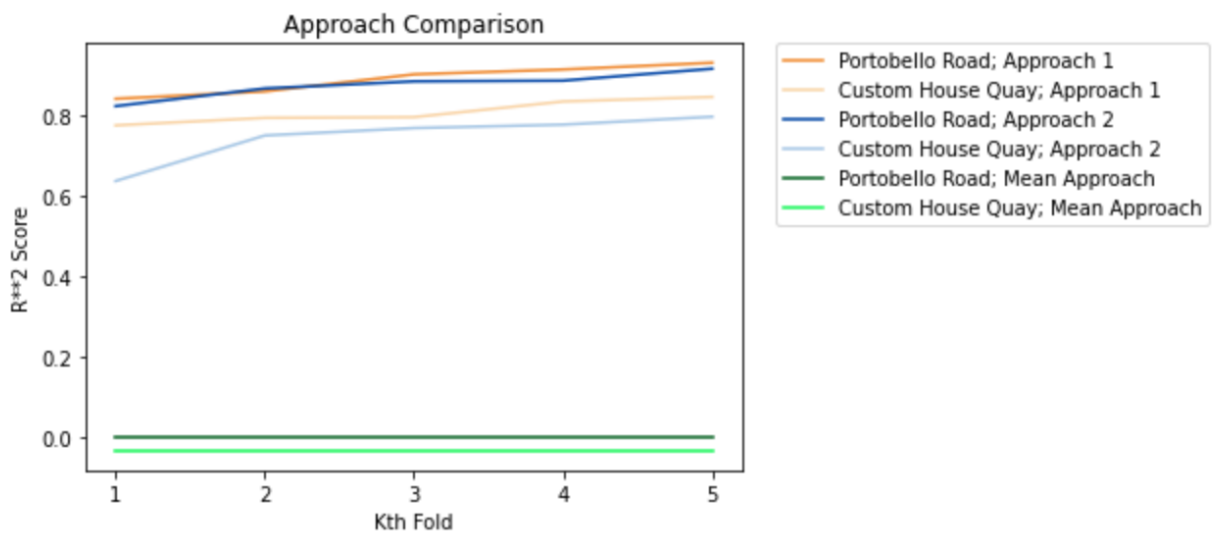


Figure 4: % bike occupancy for every station over the total duration